

# Applied Data Science with R: Data Transformation

---

Matthias Haber

27 February 2019

## **Last Weeks Homework**

---

## Quiz question: Symptoms of messy data

Which of the following is not a symptom of messy data?

- Multiple values are stored in one column

## Tidy up dataset: Read data

```
library(readxl)
library(tidyverse)
url <- paste0("http://s3.amazonaws.com/",
              "assets.datacamp.com/",
              "production/course_1294/datasets/mbta.xlsx")
download.file(url, "mbta.xlsx")
mbta <- read_excel("mbta.xlsx", skip = 1,
                  range = cell_cols(2:60))
```

# Tidy up dataset: Look at data

```
head(mbta)
```

```
## # A tibble: 6 x 59
```

```
##   mode `2007-01` `2007-02` `2007-03` `2007-04` `2007-05` `2007-06`
```

```
##   <chr> <chr>      <chr>          <dbl> <chr>      <chr>          <dbl>
```

```
## 1 All ~ NA          NA          1188. NA          NA          1246.
```

```
## 2 Boat  4           3.6           40  4.3          4.9           5.8
```

```
## 3 Bus   335.819      338.675        340. 352.162      354.367        351.
```

```
## 4 Comm~ 142.2       138.5          138. 139.5         139            143
```

```
## 5 Heav~ 435.294      448.271        459. 472.201       474.579        477.
```

```
## 6 Ligh~ 227.231       240.262        241. 255.557       248.262        246.
```

```
## # ... with 52 more variables: `2007-07` <chr>, `2007-08` <chr>,
```

```
## #   `2007-09` <dbl>, `2007-10` <chr>, `2007-11` <chr>, `2007-12` <dbl>
```

```
## #   `2008-01` <chr>, `2008-02` <chr>, `2008-03` <dbl>, `2008-04` <chr>
```

```
## #   `2008-05` <chr>, `2008-06` <dbl>, `2008-07` <chr>, `2008-08` <chr>
```

```
## #   `2008-09` <dbl>, `2008-10` <chr>, `2008-11` <chr>, `2008-12` <dbl>
```

```
## #   `2009-01` <chr>, `2009-02` <chr>, `2009-03` <dbl>, `2009-04` <chr>
```

```
## #   `2009-05` <chr>, `2009-06` <dbl>, `2009-07` <chr>, `2009-08` <chr>
```

```
## #   `2009-09` <dbl>, `2009-10` <chr>, `2009-11` <chr>, `2009-12` <dbl>
```

## Tidy up dataset: Gather years

```
mbta_tidy <- mbta %>%  
  tidyr::gather(`2007-01`:`2011-10`,  
                key = year, value = passengers,  
                convert = TRUE)
```

```
mbta_tidy
```

```
## # A tibble: 638 x 3  
##   mode          year  passengers  
##   <chr>         <chr>    <chr>  
## 1 All Modes by Qtr 2007-01 NA  
## 2 Boat            2007-01 4  
## 3 Bus              2007-01 335.819  
## 4 Commuter Rail    2007-01 142.2  
## 5 Heavy Rail       2007-01 435.294  
## 6 Light Rail       2007-01 227.231  
## 7 Pct Chg / Yr     2007-01 0.02  
## 8 Private Bus      2007-01 4.772  
## 9 RIDE             2007-01 4.9
```

## Tidy up dataset: Separate year

```
mbta_tidy <- mbta_tidy %>%  
  tidyr::separate(year, into = c("year", "month"))  
mbta_tidy
```

```
## # A tibble: 638 x 4  
##   mode          year month passengers  
##   <chr>        <chr> <chr> <chr>  
## 1 All Modes by Qtr 2007 01      NA  
## 2 Boat            2007 01      4  
## 3 Bus             2007 01    335.819  
## 4 Commuter Rail    2007 01    142.2  
## 5 Heavy Rail       2007 01    435.294  
## 6 Light Rail       2007 01    227.231  
## 7 Pct Chg / Yr     2007 01     0.02  
## 8 Private Bus      2007 01     4.772  
## 9 RIDE             2007 01     4.9  
## 10 Trackless Trolley 2007 01    12.757  
## # ... with 628 more rows
```

## Tidy up dataset: Spread mode of transportation

```
mbta_tidy <- mbta_tidy %>%  
  tidyr::spread(mode, passengers)  
mbta_tidy
```

```
## # A tibble: 58 x 13
```

```
##   year  month `All Modes by Q~ Boat  Bus   `Commuter Rail` `Heavy R  
##   <chr> <chr> <chr>          <chr> <chr> <chr>          <chr>  
## 1 2007   01    NA              4     335.~ 142.2          435.294  
## 2 2007   02    NA              3.6   338.~ 138.5          448.271  
## 3 2007   03    1187.653         40     339.~ 137.7          458.583  
## 4 2007   04    NA              4.3   352.~ 139.5          472.201  
## 5 2007   05    NA              4.9   354.~ 139          474.579  
## 6 2007   06    1245.959         5.8   350.~ 143          477.032  
## 7 2007   07    NA              6.521 357.~ 142.391          471.735  
## 8 2007   08    NA              6.572 355.~ 142.364          461.605  
## 9 2007   09    1256.571         5.469 372.~ 143.051          499.566  
## 10 2007  10    NA              5.145 368.~ 146.542          457.741  
## # ... with 48 more rows, and 6 more variables: `Light Rail` <chr>,8 `
```



## Tidy up dataset: Keep wanted columns

```
mbta_tidy <- mbta_tidy %>%  
  .[,c(1:2,6:8)]  
mbta_tidy
```

```
## # A tibble: 58 x 5  
##   year month `Commuter Rail` `Heavy Rail` `Light Rail`  
##   <chr> <chr> <chr>          <chr>          <chr>  
## 1 2007 01    142.2          435.294        227.231  
## 2 2007 02    138.5          448.271        240.262  
## 3 2007 03    137.7          458.583        241.444  
## 4 2007 04    139.5          472.201        255.557  
## 5 2007 05    139           474.579        248.262  
## 6 2007 06    143           477.032        246.108  
## 7 2007 07    142.391        471.735        243.286  
## 8 2007 08    142.364        461.605        234.907  
## 9 2007 09    143.051        499.566        265.748  
## 10 2007 10    146.542        457.741        241.434  
## # ... with 48 more rows
```

## Tidy up dataset: Gather rail modes

```
mbta_tidy <- mbta_tidy %>%  
  tidyr::gather(`Commuter Rail`, `Light Rail`,  
                key="rail_type", value = passengers)
```

```
mbta_tidy
```

```
## # A tibble: 174 x 4
```

```
##   year month rail_type    passengers  
##   <chr> <chr> <chr>         <chr>  
## 1 2007  01   Commuter Rail 142.2  
## 2 2007  02   Commuter Rail 138.5  
## 3 2007  03   Commuter Rail 137.7  
## 4 2007  04   Commuter Rail 139.5  
## 5 2007  05   Commuter Rail 139  
## 6 2007  06   Commuter Rail 143  
## 7 2007  07   Commuter Rail 142.391  
## 8 2007  08   Commuter Rail 142.364  
## 9 2007  09   Commuter Rail 143.051  
## 10 2007 10   Commuter Rail 146.542
```

## Tidy up dataset: Compute sum

```
mbta_tidy <- mbta_tidy %>%  
  dplyr::mutate(passengers = as.numeric(passengers)) %>%  
  dplyr::summarise(sum(passengers))  
mbta_tidy
```

```
## # A tibble: 1 x 1  
##   `sum(passengers)`  
##               <dbl>  
## 1             49859.
```

## Week 4: Data transformation

---

## Prerequisites

---

```
library(tidyverse)
```

# Data

336,776 flights that departed from New York City in 2013

```
# install.packages("nycflights13")  
library(nycflights13)
```

year	month	day	dep_time	sched_dep_time	dep_delay
2013	1	1	517	515	2
2013	1	1	533	529	4
2013	1	1	542	540	2
2013	1	1	544	545	-1

## Data transformation with dplyr

---



# Variable types

- `int`: integers
- `dbl`: doubles, or real numbers
- `chr`: character vectors, or strings
- `dtm`: date-times (a date + a time)
- `lgl`: logical, vectors that contain only TRUE or FALSE
- `fctr`: factors
- `date`: dates

## dplyr core functions

- `filter()`: select rows by their values
- `arrange()`: order rows
- `select()`: select columns by their names
- `mutate()`: create new variables
- `summarize()`: collapse many values down to a single summary
- `group_by()`: operate on it group-by-group
- `rename()`: rename columns
- `distinct()`: find distinct rows

Command structure (for all dplyr verbs):

- first argument is a data frame
- return value is a data frame
- nothing is modified in place

## filter()

`filter()` allows to subset observations based on their values. The function takes logical expressions and returns the rows for which all are TRUE.

### Subset Observations (Rows)



## filter()

Let's select all flights on January 1st:

```
filter(flights, month == 1, day == 1)
```

year	month	day	dep_time	sched_dep_time	dep_delay
2013	1	1	517	515	2
2013	1	1	533	529	4
2013	1	1	542	540	2
2013	1	1	544	545	-1
2013	1	1	554	600	-6
2013	1	1	554	558	-4

## filter()


`filter()` revolves around using comparison operators: `>`, `>=`, `<`, `<=`, `!=` (not equal), and `==` (equal).

`dplyr` functions like `filter()` never modify inputs but instead return a new data frame that needs to be assigned to an object if you want to save the result.

```
jan1 <- filter(flights, month == 1, day == 1)
```

## Boolean operators

`filter()` also supports the Boolean operators `&` (“and”), `|` (“or”), `!` (is “not”), and `xor` (exclusive “or”).



De Morgan's law:  $!(x \& y) = !x \mid !y$  and  $!(x \mid y) = !x \& !y$

# Boolean operators

Why does this not work?

```
filter(flights, month == 11 | 12)
```

Generally a good idea to use `x %in% y`, which will select every row where `x` is part of the values of `y`.

```
filter(flights, month %in% c(11, 12))
```

## between condition

Another useful dplyr filtering helper is `between()`. `between(x, left, right)` is equivalent to `x >= left & x <= right`.

To `filter()` all flights that departed between midnight and 6am (inclusive):

```
filter(flights, between(dep_time, 0, 600))
```



## filter() exclusion

filter() by default excludes FALSE and NA values.

```
df <- tibble(x = c(1, NA, 3))  
filter(df, x > 1)
```

```
## # A tibble: 1 x 1  
##       x  
##   <dbl>  
## 1     3
```

## filter() exclusion

If you want to preserve missing values, you have to explicitly state it.

```
filter(df, is.na(x) | x > 1)
```

```
## # A tibble: 2 x 1
```

```
##       x
```

```
##   <dbl>
```

```
## 1    NA
```

```
## 2     3
```

1. Find all flights that
  - 1.1 Had an arrival delay of two or more hours.
  - 1.2 Arrived more than two hours late, but didn't leave late.
  - 1.3 Flew to Houston (IAH or HOU).
  - 1.4 Were operated by United, American, or Delta.
  - 1.5 Departed in summer (July, August, and September).

## arrange()

`arrange()` takes a data frame and a set of column names to order the rows by. Multiple column names are evaluated subsequently.

```
arrange(flights, year, month, day)
```

year	month	day	dep_time	sched_dep_time	dep_delay
2013	1	1	517	515	2
2013	1	1	533	529	4
2013	1	1	542	540	2
2013	1	1	544	545	-1
2013	1	1	554	600	-6
2013	1	1	554	558	-4

## arrange() in descending order

By default `arrange()` sorts values in ascending order. Use `desc()` to re-order by a column in descending order.

```
arrange(flights, desc(arr_delay))
```

year	month	day	dep_time	sched_dep_time	dep_delay
2013	1	9	641	900	1301
2013	6	15	1432	1935	1137
2013	1	10	1121	1635	1126
2013	9	20	1139	1845	1014
2013	7	22	845	1600	1005
2013	4	10	1100	1900	960


### 2. Sort flights to

- 2.1 find the flight that departed the earliest (earlier than scheduled)
- 2.2 find the most delayed flight.
- 2.3 find the flight that travelled the longest and that travelled the shortest distance.

```
select()
```

`select()` is used to select a subset of variables from a dataset.

## Subset Variables (Columns)



```
select(flights, year, month, day)
```

year	month	day
2013	1	1
2013	1	1
2013	1	1
2013	1	1

## `select()`

`select()` has various helper functions:

- `everything()`: selects all variables.
- `starts_with("abc")`: matches names that begin with "abc".
- `ends_with("xyz")`: matches names that end with "xyz".
- `contains("ijk")`: matches names that contain "ijk".
- `matches("(.)\\1")`: selects variables that match a regular expression.
- `num_range("x", 1:3)` matches `x1`, `x2` and `x3`.

See `?select` for more details.



## select()

You can use `select()` to rename variables

```
select(flights, tail_num = tailnum)
```

which will drop all of the variables not explicitly mentioned.

Therefore it's better to use `rename()` instead:

```
rename(flights, tail_num = tailnum)
```

3.1 What are three distinct ways to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

3.2 What does the `one_of()` function do? Why might it be helpful in conjunction with this vector?

```
vars <- c("year", "month", "day", "dep_delay",  
          "arr_delay")
```

`mutate()` allows to add new columns to the end of your dataset that are functions of existing columns.

## Make New Variables



## mutate()

```
flights %>%  
  select(ends_with("delay"), distance, air_time) %>%  
  mutate(gain = arr_delay - dep_delay,  
         speed = distance / air_time * 60  
  )
```

dep_delay	arr_delay	distance	air_time	gain	speed
2	11	1400	227	9	370.0441
4	20	1416	227	16	374.2731
2	33	1089	160	31	408.3750
-1	-18	1576	183	-17	516.7213
-6	-25	762	116	-19	394.1379
-4	12	719	150	16	287.6000

## transmute()

Use `transmute()` to only keep the new variables:

```
transmute(flights,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours  
)
```

## Functions to use with `mutate()`

There are many functions for creating new variables with `mutate()`:

- Arithmetic operators: `+`, `-`, `*`, `/`, `^` (e.g. `air_time / 60`).
- Aggregate functions: `sum(x)` `mean(y)` (e.g. `mean(dep_delay)`).
- Modular arithmetic: `%/%` (integer division) and `%%` (remainder), where `x == y * (x %/% y) + (x %% y)`.
- Logs: `log()`, `log2()`, `log10()`.
- Offsets: `lead()` and `lag()` (e.g. `x - lag(x)`).
- Cumulative and rolling aggregates: `cumsum()`, `cumprod()`, `cummin()`, `cummax()`, `cummean()`.
- Logical comparisons, `<`, `<=`, `>`, `>=`, `!=`.
- Ranking: `min_rank()`, `row_number()`, `dense_rank()`, `percent_rank()`, `cume_dist()`, `ntile()`.

4. Use `mutate()` to
  - 4.1 Create new variables for `dep_time` and `sched_dep_time` that measure time in the number of minutes after midnight.
  - 4.2 Compare `air_time` with `arr_time - dep_time`. What do you see? What do you need to do to fix it?
  - 4.3 Find the 10 most delayed flights using a ranking function. How do you want to handle ties?

## summarize()

summarize() collapses a data frame to a single row.

### Summarise Data



```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1
```

```
##   delay
```


```
##   <dbl>
```

```
## 1  12.6
```



## `summarize()` with `group_by()`

`summarize()` is most effectively used with `group_by()`, which changes the unit of analysis from the complete dataset to individual groups.



Grouping is most useful in conjunction with `summarise()`, but you can also do convenient operations with `mutate()` and `filter()`.

## summarize() with group\_by()

For example, to get the average delay per date

```
flights %>%  
  group_by(year, month, day) %>%  
  summarise(delay = mean(dep_delay, na.rm = TRUE))
```

## `summarize()` `count`

For aggregations it is generally a good idea to include a count `n()`. For example, let's look at the (not cancelled) planes that have the highest average delays:

```
flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))  
  group_by(tailnum) %>%  
  summarise(delay = mean(arr_delay)) %>%  
  arrange(delay)
```

There are a number of useful summary functions:

- Measures of location: `mean(x)`, `sum(x)`, `median(x)`.
- Measures of spread: `sd(x)`, `IQR(x)`, `mad(x)`.
- Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`.
- Measures of position: `first(x)`, `nth(x, 2)`, `last(x)`.
- Counts: `n()`, `sum(!is.na(x))`, `n_distinct(x)`.
- Counts and proportions of logical values: `sum(x > 10)`, `mean(y == 0)`.

5. Use `summarize()` to
  - 5.1 Look at the number of cancelled flights per day. Is there a pattern? Is the proportion of cancelled flights related to the average delay?
  - 5.2 Find the carrier with the worst delays.

# Homework Exercises

---

# Homework Exercises

For this week's homework exercises go to Moodle and answer the Quiz posted in the Week 4: Data Transformation section.

Deadline: Tuesday, March 5.

**That's it for today. Questions?**