

Applied Data Science with R: Web scraping

Matthias Haber

27 March 2019

Prerequisites

Packages

```
library(tidyverse)
library(stringr)
library(rvest) #install.packages("rvest")
library(RCurl) #install.packages("RCurl")
library(httr) #install.packages("httr")
library(httpuv) #install.packages("httpuv")
library(rtweet) #install.packages("tweet")
library(RSelenium) #install.packages("RSelenium")
```

Last week's homework

Extract name from email

```
email <- c("tom@hogwarts.com",  
           "tom.riddle@hogwarts.com",  
           "tom@hogwarts.eu.com",  
           "potter@hogwarts.com",  
           "harry@hogwarts.com",  
           "hermione+witch@hogwarts.com")  
str_extract(email, "^([\\w\\.\\+]+)*")
```

```
## [1] "tom"           "tom.riddle"    "tom"           "  
## [5] "harry"         "hermione+witch"
```

extract image files

```
files <- c(".bash_profile",  
  "workspace.doc",  
  "img0912.jpg",  
  "updated_img0912.png",  
  "documentation.html",  
  "favicon.gif",  
  "img0912.jpg.tmp",  
  "access.lock")  
str_extract(files, ".*\\.(jpg|png|gif)$")
```

```
## [1] NA NA "img0912.jpg"  
## [4] "updated_img0912.png" NA "favicon.gif"  
## [7] NA NA
```

Extract line number from log files

```
log <- c("W/dalvikvm( 1553): threadid=1: uncaught exception",  
        "E/( 1553): FATAL EXCEPTION: main",  
        "E/( 1553): java.lang.StringIndexOutOfBoundsException",  
        "E/( 1553):   at widget.List.makeView(ListView.java:1727)",  
        "E/( 1553):   at widget.List.fillDown(ListView.java:652)",  
        "E/( 1553):   at widget.List.fillFrom(ListView.java:709)")  
str_extract(log, "(?<=java:).+?(?=\\\\))")
```

```
## [1] NA      NA      NA      "1727" "652"  "709"
```

RStudio string cheatsheet

Webscraping intro

Example from my own work

My task was to find out how people talk about the Mercedes-Benz GLS on a forum: mbworld.org

Browsing

- you click on something
- browser sends request to server that hosts website
- server returns resource (often an HTML document)
- browser interprets HTML and renders it in a nice fashion

Scraping with R

- you manually specify a resource
- R sends request to server that hosts website
- server returns resource
- R parses HTML (i.e., interprets the structure), but does not render it in a nice fashion
- it's up to you to tell R which parts of the structure to focus on and what content to extract

- **web pages** (e.g. <http://example.com>)
- **web formats** (XML, HTML, JSON, ...)
- **web frameworks** (HTTP, URL, APIs, ...)
- **social media** (Twitter, Facebook, LinkedIn, Snapchat, Tumblr, ...)
- **data in the web** (speeches, laws, policy reports, news, ...)
- **web data** (page views, page ranks, IP-addresses, ...)

Before scraping, do some googling!

- If the resource is well-known, someone else has probably built a tool which solves the problem for you.
- ropensci has a ton of R packages providing easy-to-use interfaces to open data.
- The Web Technologies and Services CRAN Task View is a great overview of various tools for working with data that lives on the web in R.

What's HTML?

HyperText Markup Language

- markup language = plain text + markups
- standard for the construction of websites
- relevance for web scraping: web architecture is important because it determines where and how information is stored

Inspect the source code in your browser

Firefox 1. right click on page 2. select “View Page Source”

Chrome 1. right click on page 2. select “View page source”

Safari 1. click on “Safari” 2. select “Preferences” 3. go to “Advanced” 4. check “Show Develop menu in menu bar” 5. click on “Develop” 6. select “Show Page Source.”

Cascading Style Sheets

- style sheet language to give browsers information of how to render HTML documents
- CSS code can be stored within an HTML document or in an external CSS file
- selectors, i.e. patterns used to specify which elements to format in a certain way, can be used to address the elements we want to extract information from
- works via tag name (e.g., `<h2>`, `<p>`) or element attributes `id` and `class`

How does it work?

CSS Diner

- XPath is a query language for selecting nodes from an XML-style document (including HTML)
- provides just another way of extracting data from static webpages
- you can also use XPath with R, it can be more powerful than CSS selectors



WIKIPEDIA
The Free Encyclopedia

- Main page
- Contents
- Featured content
- Current events
- Random article
- Donate to Wikipedia
- Wikipedia store

Interaction

- Help
- About Wikipedia
- Community portal
- Recent changes
- Contact page

Tools

Not logged in [Talk](#) [Contributions](#) [Create account](#) [Log in](#)

Article [Talk](#)

[Read](#) [Edit](#) [View history](#)

Table (information)

From Wikipedia, the free encyclopedia

"Tabular" redirects here. For the typewriter key, see [tab key](#).

For sortable tables in Wikipedia, see [Help:Sorting](#)



This article may **require cleanup** to meet Wikipedia's [quality standards](#). No [cleanup reason](#) has been specified. Please help [improve this article](#) if you can. *(October 2009)* ([Learn how and when to remove this template message](#))

A **table** is an arrangement of [data](#) in rows and columns, or possibly in a more complex structure. Tables are widely used in [communication](#), [research](#), and [data analysis](#). Tables appear in print media, handwritten notes, computer software, architectural ornamentation, traffic signs, and many other places. The precise conventions and terminology for describing tables vary depending on the context. Further, tables differ significantly in variety, structure, flexibility, notation, representation and use.^{[1][2][3][4][5]} In books and technical articles, tables are typically presented apart from the main text in numbered and captioned [floating blocks](#).

Purchased Equipments (June, 2006)			
Item Num#	Item Picture	Item Description	Price
		Shipping Handling, Installation, etc	Expense
1.		IBM Case Computer	\$ 400.00
		Shipping Handling, Installation, etc	\$ 20.00
2.		1GB RAM Module for Computer	\$ 50.00
		Shipping Handling, Installation, etc	\$ 14.00
Purchased Equipments (June, 2006)			

Inspecting elements

Simple table [\[edit\]](#)

The following illustrates a simple table with three columns and six rows. display the column names. This is traditionally called a "header row".

Age table

First name	Last name	Age	← → ↻ ☆
Bielat	Adamczak	24	Save Page As...
Blaszczyk	Kostrzewski	25	View Background Image Select All This Frame ▶
Olatunkboh	Chijiaku	22	View Page Source View Page Info
Adrienne	Anthoula	22	Inspect Element
Axelia	Athanasios	22	Inspect Element with Firebug
Jon-Kabat	Zinn	22	1Password

Multi-dimensional table

Hover to find desired elements

Runa Simi
Simple English
Slovenčina
Svenska
தமிழ்
ไทย
Українська

עברית
中文
Edit links

Simple table [\[edit\]](#)

The following illustrates a simple table with three columns and six rows. The first row is not counted, because it is only used to display the column names. This is traditionally called a "header row".

Table.wikitable 253 x 245

First name	Last name	Age
Bielat	Adamczak	24
Błaszczuk	Kostrzewski	25
Olatunkboh	Chijiaku	22
Adrienne	Anthoula	22
Axelia	Athanasios	22
Jon-Kabat	Zinn	22

Multi-dimensional table [\[edit\]](#)

The concept of **dimension** is also a part of basic terminology.^[7] Any "simple" table can

Wikicontent.mw-body > divbodyContent.mw-body-content > divmw-content-text.mw-content-tr > **table.wikitable** > tbody > tr > td > Q

```
<table class="wikitable" style="text-align:center">
  <caption></caption>
  <tbody>
    <tr></tr>
    <tr>
      <td></td>
      <td></td>
      <td></td>
    </tr>
    <tr></tr>
    <tr></tr>
    <tr></tr>
  </tbody>
</table>
```

Rules Computed Fonts Box K

```
element {
  text-align: center;
}
table.wikitable {
  margin: 1em 0px;
  background-color: #F9F9F9;
  border: 1px solid #AAA;
  border-collapse: collapse;
  color: #000;
}
table {
```

rvest is a nice R package for web-scraping by (you guessed it) Hadley Wickham.

- see also: <https://github.com/hadley/rvest>
- convenient package to scrape information from web pages
- builds on other packages, such as xml2 and httr
- provides very intuitive functions to import and process webpages

Basic workflow of scraping with rvest

1. specify URL

```
"http://en.wikipedia.org/wiki/Table_(information)" %>%
```

2. download static HTML behind the URL and parse it into an XML file

```
read_html() %>%
```

3. extract specific nodes with CSS (or XPath)

```
html_node(".wikitable") %>%
```

4. extract content from nodes

```
html_table(fill = TRUE)
```

##	First name	Last name	Age
## 1	Tinu	Elejogun	14
## 2	Blaszczyk	Kostrzewski	25
## 3	Lily	McGarrett	18
## 4	Olatunkbo	Chijiaku	22
## 5	Adrienne	Anthoula	22
## 6	Amelia	Athenagias	22

- Selectorgadget is a Chrome browser extension for quickly extracting desired parts of an HTML page.
- to learn about it, use `vignette("selectorgadget")`
- to install it, visit <http://selectorgadget.com/>

```
url <- "http://spiegel.de/schlagzeilen"  
css <- ".schlagzeilen-headline"  
url_parsed <- read_html(url)  
html_nodes(url_parsed, css = css) %>% html_text
```

Technologies and Packages

- Regular Expressions / String Handling
 - **stringr**, stringi
- HTML / XML / XPath / CSS Selectors
 - **rvest**, xml2, XML
- JSON
 - **jsonlite**, RJSONIO, rjson
- HTTP / HTTPS
 - **httr**, curl, RCurl
- Javascript / Browser Automation
 - **RSelenium**
- URL
 - **urltools**

APIs

- Application programming interfaces (APIs) allow a direct connection to data hosted on a website
- Most modern APIs use HTTP (HyperText Transfer Protocol) for communication and data transfer between server and client
- R package `httr` as a good-to-use HTTP client interface
- Most web data APIs return data in JSON or XML format
- R packages `jsonlite` and `xml2` good to process JSON or XML-style data

- If you want to tap an existing API, you have to
 - figure out how it works (what requests/actions are possible, what endpoints exist)
 - (register to use the API)
 - formulate queries to the API from within R
 - process the incoming data

Twitter has two types of APIs

- REST APIs → reading/writing/following/etc.
- Streaming APIs → low latency access to 1% of global stream - public, user and site streams
- authentication via OAuth
- documentation at <https://dev.twitter.com/overview/documentation>

Accessing the Twitter APIs

To access the REST and streaming APIs, all you need is a Twitter account and you can be up in running in minutes!

Simply send a request to Twitter's API (with a function like `search_tweets()`) during an interactive session of R, authorize the embedded `rstats2twitter` app (approve the browser popup), and your token will be created and saved/stored (for future sessions) for you!

You can obtain a developer account to get more stability and permissions.

Use twitter in R

```
library(rtweet)

## search for 1000 tweets using the #niewiedercdu hashtag
tweets <- search_tweets(
  "#nieMehrCDU", n = 1000, include_rts = FALSE
)
```

Find out what else you can do with the rtweet package:

<https://github.com/mkearney/rtweet>

Scraping dynamic pages

To scrape dynamic web pages (e.g. with Java running in the background) use RSelenium

Group excercises

Breweries in Germany

Fetch a list of cities with breweries in Germany:

<http://www.biermap24.de/brauereiliste.php>

Repeat playing CSS diner and complete all levels!

Homework Exercises

Homework Exercises

For this week's homework exercises go to Moodle and complete the assignment posted in the Web Scraping section.

Deadline: Tuesday, April 2.

That's it for today. Questions?